

kwd_project_code

February 3, 2020

0.1 Przygotowanie środowiska

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import category_encoders as ce
import sklearn as skl
```

```
[2]: np.__version__
```

```
[2]: '1.17.2'
```

0.2 Załadowanie danych z pliku

```
[3]: data = pd.read_csv('./online_video_dataset/transcoding_mesurment.tsv',
    ↪delimiter='\t')
data.head(2)
```

```
[3]:
```

	id	duration	codec	width	height	bitrate	framerate	i	p	\
0	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
1	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	

	b	...	p_size	b_size	size	o_codec	o_bitrate	o_framerate	o_width	\
0	0	...	825054	0	889537	mpeg4	56000	12.0	176	
1	0	...	825054	0	889537	mpeg4	56000	12.0	320	

	o_height	umem	utime
0	144	22508	0.612
1	240	25164	0.980

[2 rows x 22 columns]

0.3 Wstępna analiza danych

```
[4]: print(data.shape)
print(data.columns)
```

```
(68784, 22)
Index(['id', 'duration', 'codec', 'width', 'height', 'bitrate', 'framerate',
      'i', 'p', 'b', 'frames', 'i_size', 'p_size', 'b_size', 'size',
      'o_codec', 'o_bitrate', 'o_framerate', 'o_width', 'o_height', 'umem',
      'utime'],
      dtype='object')
```

Opis danych: - id = Youtube videp id - duration = duration of video - bitrate bitrate(video) = video bitrate - height = height of video in pixles - width = width of video in pixles - frame rate = actual video frame rate - frame rate(est.) = estimated video frame rate - codec = coding standard used for the video - category = YouTube video category - url = direct link to video (has expiration date) - i = number of i frames in the video (complete image) - p = number of p frames in the video (predicted picture) - b = number of b frames in the video (bidirectional predicted picture) - frames = number of frames in video - i_size = total size in byte of i videos - p_size = total size in byte of p videos - b_size = total size in byte of b videos - size = total size of video - o_codec = output codec used for transcoding - o_bitrate = output bitrate used for transcoding - o_framerate = output framerate used for transcoding - o_width = output width in pixel used for transcoding - o_height = output height used in pixel for transcoding - umem = total codec allocated memory for transcoding - utime = total transcoding time for transcoding

0.4 Analiza eksploracyjna

```
[5]: # typy danych w kolumnach
      data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 68784 entries, 0 to 68783
Data columns (total 22 columns):
id                68784 non-null object
duration          68784 non-null float64
codec             68784 non-null object
width             68784 non-null int64
height            68784 non-null int64
bitrate           68784 non-null int64
framerate         68784 non-null float64
i                 68784 non-null int64
p                 68784 non-null int64
b                 68784 non-null int64
frames            68784 non-null int64
i_size            68784 non-null int64
p_size            68784 non-null int64
b_size            68784 non-null int64
size              68784 non-null int64
o_codec           68784 non-null object
o_bitrate         68784 non-null int64
o_framerate       68784 non-null float64
o_width           68784 non-null int64
o_height          68784 non-null int64
```

```
umem          68784 non-null int64
utime         68784 non-null float64
dtypes: float64(4), int64(15), object(3)
memory usage: 11.5+ MB
```

```
[6]: # Jakie są podstawowe metryki statystyczne poszczególnych kolumn?
data.describe()
```

```
[6]:
```

	duration	width	height	bitrate	framerate \
count	68784.000000	68784.000000	68784.000000	6.878400e+04	68784.000000
mean	286.413921	624.934171	412.572226	6.937015e+05	23.241321
std	287.257650	463.169069	240.615472	1.095628e+06	7.224848
min	31.080000	176.000000	144.000000	8.384000e+03	5.705752
25%	106.765000	320.000000	240.000000	1.343340e+05	15.000000
50%	239.141660	480.000000	360.000000	2.911500e+05	25.021740
75%	379.320000	640.000000	480.000000	6.529670e+05	29.000000
max	25844.086000	1920.000000	1080.000000	7.628466e+06	48.000000

	i	p	b	frames	i_size \
count	68784.000000	68784.000000	68784.000000	68784.000000	6.878400e+04
mean	100.868312	6531.692210	9.147854	6641.708377	2.838987e+06
std	84.764791	6075.871744	92.516177	6153.342453	4.325137e+06
min	7.000000	175.000000	0.000000	192.000000	1.164800e+04
25%	39.000000	2374.000000	0.000000	2417.000000	3.933950e+05
50%	80.000000	5515.000000	0.000000	5628.000000	9.458650e+05
75%	138.000000	9155.000000	0.000000	9232.000000	3.392479e+06
max	5170.000000	304959.000000	9407.000000	310129.000000	9.082855e+07

	p_size	b_size	size	o_bitrate	o_framerate \
count	6.878400e+04	68784.0	6.878400e+04	6.878400e+04	68784.000000
mean	2.218057e+07	0.0	2.502294e+07	1.395036e+06	21.190862
std	5.097306e+07	0.0	5.414402e+07	1.749352e+06	6.668703
min	3.384500e+04	0.0	1.918790e+05	5.600000e+04	12.000000
25%	1.851539e+06	0.0	2.258222e+06	1.090000e+05	15.000000
50%	6.166260e+06	0.0	7.881069e+06	5.390000e+05	24.000000
75%	1.515506e+07	0.0	1.977335e+07	3.000000e+06	25.000000
max	7.689970e+08	0.0	8.067111e+08	5.000000e+06	29.970000

	o_width	o_height	umem	utime
count	68784.000000	68784.000000	68784.000000	68784.000000
mean	802.336357	503.825541	228224.717900	9.996355
std	609.959797	315.970438	97430.878373	16.107429
min	176.000000	144.000000	22508.000000	0.184000
25%	320.000000	240.000000	216820.000000	2.096000
50%	480.000000	360.000000	219480.000000	4.408000
75%	1280.000000	720.000000	219656.000000	10.433000
max	1920.000000	1080.000000	711824.000000	224.574000

```
[7]: # czy są zduplikowane wartości?  
len(data[data.duplicated()])
```

[7]: 0

0.4.1 Korelacja

Na co zwrócić uwagę? * zostawić zmienne, skorelowane z targetem; * usunąć zmienne skorelowane ze sobą.

Informacja przydaje się na etapie **Feature Selection** - co chcemy włączyć w dataset treningowy.

```
[8]: # sns.set()  
  
# corr = data.sample(1000).corr()  
  
# fig, ax = plt.subplots(figsize=(50,30)) # Sample figsize in inches  
# sns.heatmap(corr, annot=True, linewidths=.5, ax=ax)  
# plt.show()
```

Widoczna jest silna korelacja pomiędzy wymiarami docelowego pliku a czasem transkodowania.


0.4.2 Analiza pól numerycznych

```
[9]: numeric_columns = data.select_dtypes(include=np.number).columns  
numeric_columns
```


```
[9]: Index(['duration', 'width', 'height', 'bitrate', 'framerate', 'i', 'p', 'b',  
        'frames', 'i_size', 'p_size', 'b_size', 'size', 'o_bitrate',  
        'o_framerate', 'o_width', 'o_height', 'umem', 'utime'],  
        dtype='object')
```

```
[10]: for col in numeric_columns:  
    print(f'kolumna: {col}')  
    print(f'unikalnych wartości: {len(data[col].unique())}')  
    data[col].plot.hist(bins=40)  
    plt.show()
```


kolumna: duration
unikalnych wartości: 1086




kolumna: width
unikalnych wartości: 6




kolumna: height
unikalnych wartości: 6




kolumna: bitrate
 unikalnych wartości: 1095




kolumna: framerate
 unikalnych wartości: 261




kolumna: i
 unikalnych wartości: 306




kolumna: p
 unikalnych wartości: 1042




kolumna: b
unikalnych wartości: 20



kolumna: frames
unikalnych wartości: 1044




kolumna: `i_size`
 unikalnych wartości: 1099




kolumna: `p_size`

unikalnych wartości: 1099




kolumna: b_size


unikalnych wartości: 1




kolumna: size
unikalnych wartości: 1099




kolumna: o_bitrate
unikalnych wartości: 7




kolumna: o_framerate
unikalnych wartości: 5




kolumna: o_width
unikalnych wartości: 6




kolumna: o_height
unikalnych wartości: 6



kolumna: umem
unikalnych wartości: 9395



kolumna: utime
unikalnych wartości: 10960



```
[11]: print(f"unikalnych wartości w 'b_size': {len(data['b_size'].unique())}")
      print(f"Wierszy z polem 'b' != 0: {len(data[data['b'] > 0])}")
      print(f"Wierszy z polem 'b' != 0: {len(data[data['b'] > 0])/len(data)*100}%")
```

```
unikalnych wartości w 'b_size': 1
Wierszy z polem 'b' != 0: 859
Wierszy z polem 'b' != 0: 1.248836938822982%
```

Wnioski Kolumna 'b_size' nie zawiera niezerowych wartości. Można ją usunąć. Kolumna 'b' jest niezerowa w ~1% przypadków.


Analiza kolumn kategorycznych

```
[12]: non_numeric_columns = data.select_dtypes(exclude=np.number).columns
      non_numeric_columns
```


```
[12]: Index(['id', 'codec', 'o_codec'], dtype='object')
```

```
[13]: for col in non_numeric_columns:
      print(f'kolumna: {col}')
      print(f'unikalnych wartości: {len(data[col].unique())}')
      if len(data[col].unique()) < 15:
          data[col].value_counts().plot.bar()
          plt.show()
```

```
kolumna: id
unikalnych wartości: 1099
kolumna: codec
unikalnych wartości: 4
```



kolumna: o_codec
unikalnych wartości: 4



0.4.3 Wartości brakujące

```
[14]: data.isnull().sum()
```

```
[14]: id          0
      duration    0
      codec       0
      width       0
      height      0
      bitrate     0
      framerate   0
      i           0
      p           0
      b           0
      frames      0
      i_size      0
      p_size      0
      b_size      0
      size        0
      o_codec     0
      o_bitrate   0
      o_framerate 0
      o_width     0
      o_height    0
      umem        0
      utime       0
      dtype: int64
```

0.5 Outliers

```
[15]: data_2 = data.copy()
```

Do obliczenia outlierów pomocniczo wykorzystujemy poniższe wzory: - IQR (interquartile range) = $P(75) - P(25)$ - Dolna_granica = $P(25) - 1,5IQR$ - Górna_granica = $P(75) + 1,5IQR$ - (P - percentyl)

```
[16]: def outliers_range(data, column_name):
      rows = data[column_name]
      iqr = np.nanpercentile(rows, 75) - np.nanpercentile(rows, 25)
      lower = (np.nanpercentile(rows, 25) - 1.5*iqr)
      upper = (np.nanpercentile(rows, 75) + 1.5*iqr)
      return lower, upper
```

0.5.1 Kolumny numeryczne

```
[17]: numeric_columns = data_2.select_dtypes(include=np.number).columns
      numeric_data = data_2[numeric_columns]
      numeric_data.head()
```

```
[17]:
```



	duration	width	height	bitrate	framerate	i	p	b	frames	i_size	\
0	130.35667	176	144	54590	12.0	27	1537	0	1564	64483	
1	130.35667	176	144	54590	12.0	27	1537	0	1564	64483	
2	130.35667	176	144	54590	12.0	27	1537	0	1564	64483	
3	130.35667	176	144	54590	12.0	27	1537	0	1564	64483	
4	130.35667	176	144	54590	12.0	27	1537	0	1564	64483	

	p_size	b_size	size	o_bitrate	o_framerate	o_width	o_height	umem	\
0	825054	0	889537	56000	12.0	176	144	22508	
1	825054	0	889537	56000	12.0	320	240	25164	
2	825054	0	889537	56000	12.0	480	360	29228	
3	825054	0	889537	56000	12.0	640	480	34316	
4	825054	0	889537	56000	12.0	1280	720	58528	

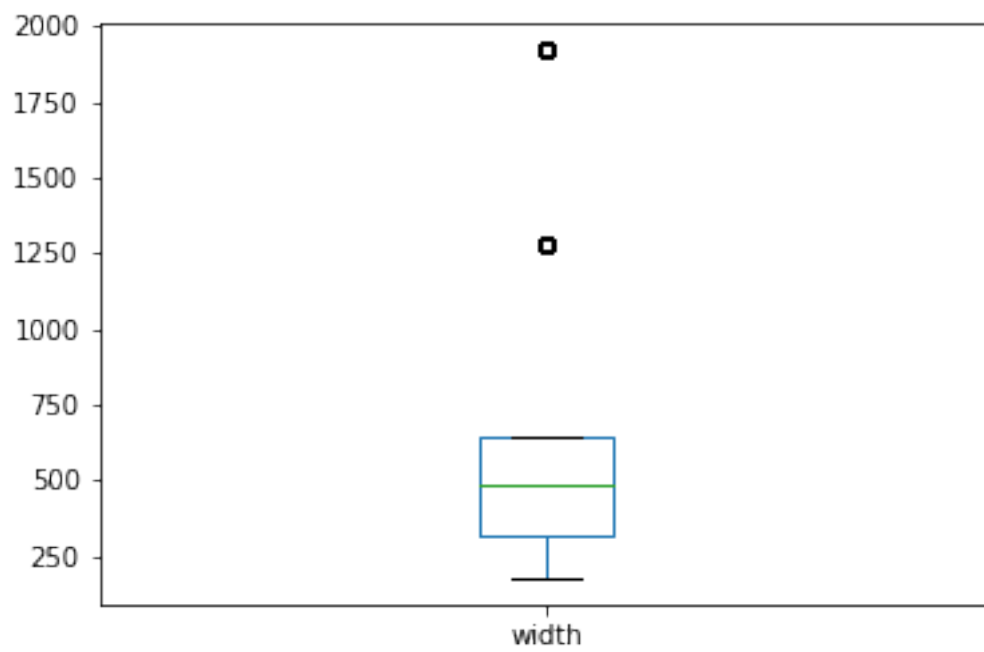

	utime
0	0.612
1	0.980
2	1.216
3	1.692
4	3.456

```
[18]: # wizualizacja rozkładu danych poszczególnych kolumn
      for column in numeric_data:
          print(f'Kolumna: {column}')
          numeric_data[column].plot.hist()
          plt.show()
          numeric_data[column].plot.box()
          plt.show()
```



Kolumna: duration





Kolumna: width



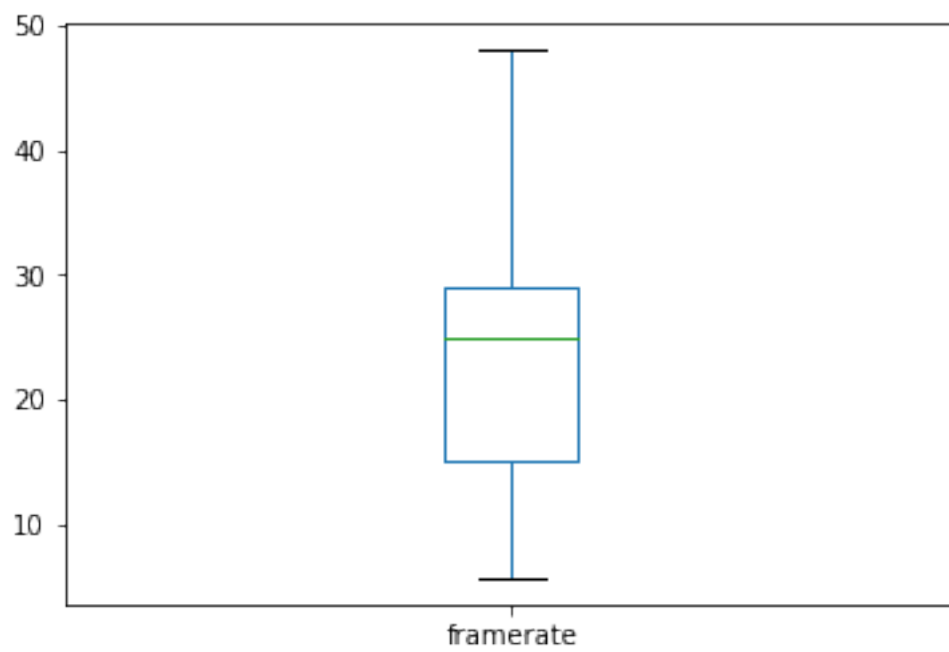

Kolumna: height





Kolumna: bitrate





Kolumna: framerate





Kolumna: i





Kolumna: p





Kolumna: b





Kolumna: frames





Kolumna: `i_size`





Kolumna: `p_size`





Kolumna: b_size





Kolumna: size





Kolumna: o_bitrate





Kolumna: o_framerate





Kolumna: o_width





Kolumna: o_height



Kolumna: umem



Kolumna: utime



Kolumna “frames”

```
[19]: lower, upper = outliers_range(numeric_data, 'frames')
      lower, upper
```

```
[19]: (-7805.5, 19454.5)
```

```
[20]: num = len(numeric_data[numeric_data["frames"]>upper])
print(f'wartości >{upper}: {num} ({round(num/len(data_2)*100, 2)}%)')
num = len(numeric_data[numeric_data["frames"]>30000])
print(f'wartości >30k: {num} ({round(num/len(data_2)*100, 2)}%)')
num = len(numeric_data[numeric_data["frames"]>40000])
print(f'wartości >35k: {num} ({round(num/len(data_2)*100, 2)}%)')
```


wartości >19454.5: 3076 (4.47%)

wartości >30k: 869 (1.26%)

wartości >35k: 20 (0.03%)

```
[21]: # Obcięcie outlierów ( usunięcie wierszy z wart. >35k)
data_2 = data_2[data_2['frames'] <= 35000 ]

data_2['frames'].plot.box()
plt.show()
```



Kolumna “duration”

```
[22]: lower, upper = outliers_range(numeric_data, 'duration')
lower, upper
```


```
[22]: (-302.0675, 788.1524999999999)
```

```
[23]: num = len(numeric_data[numeric_data["duration"]>788])
print(f'wartości >788: {num} ({round(num/len(data_2)*100, 2)}%)')
num = len(numeric_data[numeric_data["duration"]>1000])
print(f'wartości >1000: {num} ({round(num/len(data_2)*100, 2)}%)')
num = len(numeric_data[numeric_data["duration"]>1500])
print(f'wartości >1500: {num} ({round(num/len(data_2)*100, 2)}%)')
num = len(numeric_data[numeric_data["duration"]>2000])
print(f'wartości >2000: {num} ({round(num/len(data_2)*100, 2)}%)')
```

```
wartości >788: 3080 (4.48%)
wartości >1000: 1718 (2.5%)
wartości >1500: 865 (1.26%)
wartości >2000: 13 (0.02%)
```

```
[24]: # Obcięcie outlierów ( usunięcie wierszy z wart. >2k)
data_2 = data_2[data_2['duration'] <= 2000 ]

data_2['duration'].plot.box()
plt.show()
```



0.6 Feature Engineering

```
[25]: data_3 = data_2.copy()
```

```
[26]: codecs_with_b      = data_3[data_3['b']==0]['codec'].unique()
      codecs_without_b = data_3[data_3['b']!=0]['codec'].unique()
      print(f"codecs_with_b: {codecs_with_b}")
      print(f"codecs_without_b: {codecs_without_b}")
```

```
codecs_with_b: ['mpeg4' 'h264' 'vp8' 'flv']
codecs_without_b: ['h264']
```

Wniosek: Parametr 'b' jest niezerowy tylko w przypadku użycia kodeka 'h264'

Dodatkowe pole - pixels (width*height)

```
[27]: data_3['pixels'] = data_3['width'] * data_3['height']
      data_3.head()
```

```
[27]:
```

	id	duration	codec	width	height	bitrate	framerate	i	p	\
0	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
1	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
2	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
3	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
4	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	

	b	...	b_size	size	o_codec	o_bitrate	o_framerate	o_width	o_height	\
0	0	...	0	889537	mpeg4	56000	12.0	176	144	
1	0	...	0	889537	mpeg4	56000	12.0	320	240	
2	0	...	0	889537	mpeg4	56000	12.0	480	360	
3	0	...	0	889537	mpeg4	56000	12.0	640	480	
4	0	...	0	889537	mpeg4	56000	12.0	1280	720	

	umem	utime	pixels
0	22508	0.612	25344
1	25164	0.980	25344
2	29228	1.216	25344
3	34316	1.692	25344
4	58528	3.456	25344

[5 rows x 23 columns]

Dodatkowe pole - o_pixels (o_width*o_height)

```
[28]: data_3['o_pixels'] = data_3['o_width'] * data_3['o_height']
      data_3.head()
```

```
[28]:
```

	id	duration	codec	width	height	bitrate	framerate	i	p	\
0	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
1	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
2	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	
3	04t6-jw9czg	130.35667	mpeg4	176	144	54590	12.0	27	1537	


```
4  04t6-jw9czg  130.35667  mpeg4    176    144    54590        12.0  27  1537
```

```

      b  ...    size  o_codec  o_bitrate  o_framerate  o_width  o_height  umem  \
0  0  ...  889537    mpeg4    56000        12.0        176        144  22508
1  0  ...  889537    mpeg4    56000        12.0        320        240  25164
2  0  ...  889537    mpeg4    56000        12.0        480        360  29228
3  0  ...  889537    mpeg4    56000        12.0        640        480  34316
4  0  ...  889537    mpeg4    56000        12.0       1280        720  58528

```

```

      utime  pixels  o_pixels
0  0.612    25344    25344
1  0.980    25344    76800
2  1.216    25344   172800
3  1.692    25344   307200
4  3.456    25344   921600

```

[5 rows x 24 columns]

0.7 Feature selection - usunięcie zbędnych kolumn

- pole id
- pole b_size (brak niezerowych wartości)
- pole umem (wartość wynikowa, jednak nie podlegająca analizie)

```
[29]: data_4 = data_3.copy()
```

```
[30]: data_4 = data_4.drop(['id', 'b_size', 'umem'], axis=1)
      data_4.head(3)
```

```
[30]:
      duration  codec  width  height  bitrate  framerate  i    p  b  frames  \
0  130.35667  mpeg4    176    144    54590        12.0  27  1537  0    1564
1  130.35667  mpeg4    176    144    54590        12.0  27  1537  0    1564
2  130.35667  mpeg4    176    144    54590        12.0  27  1537  0    1564

```

```

      ...  p_size    size  o_codec  o_bitrate  o_framerate  o_width  o_height  \
0  ...  825054  889537    mpeg4    56000        12.0        176        144
1  ...  825054  889537    mpeg4    56000        12.0        320        240
2  ...  825054  889537    mpeg4    56000        12.0        480        360

```

```

      utime  pixels  o_pixels
0  0.612    25344    25344
1  0.980    25344    76800
2  1.216    25344   172800

```

[3 rows x 21 columns]

0.8 Data Preparation - przygotowanie danych do uczenia

- kodowanie pól katagorycznych
- skalowanie danych

```
[31]: data_5 = data_4.copy()
data_5.head()
```

```
[31]:      duration  codec  width  height  bitrate  framerate  i    p  b  frames  \
0  130.35667  mpeg4    176    144    54590         12.0  27  1537  0    1564
1  130.35667  mpeg4    176    144    54590         12.0  27  1537  0    1564
2  130.35667  mpeg4    176    144    54590         12.0  27  1537  0    1564
3  130.35667  mpeg4    176    144    54590         12.0  27  1537  0    1564
4  130.35667  mpeg4    176    144    54590         12.0  27  1537  0    1564

      ...  p_size    size  o_codec  o_bitrate  o_framerate  o_width  o_height  \
0  ...  825054  889537    mpeg4    56000         12.0        176        144
1  ...  825054  889537    mpeg4    56000         12.0        320        240
2  ...  825054  889537    mpeg4    56000         12.0        480        360
3  ...  825054  889537    mpeg4    56000         12.0        640        480
4  ...  825054  889537    mpeg4    56000         12.0       1280       720

      utime  pixels  o_pixels
0  0.612   25344    25344
1  0.980   25344    76800
2  1.216   25344   172800
3  1.692   25344   307200
4  3.456   25344   921600

[5 rows x 21 columns]
```

```
[32]: categorical_cols = ['codec', 'o_codec']
for column in categorical_cols:
    print(column)
    print(data_5[column].unique())
```

```
codec
['mpeg4' 'h264' 'vp8' 'flv']
o_codec
['mpeg4' 'vp8' 'flv' 'h264']
```

```
[33]: ohe = ce.OneHotEncoder(cols=categorical_cols, return_df=True,
        use_cat_names=True, handle_unknown=0)
data_5 = ohe.fit_transform(data_5)
data_5.head()
```

```
[33]:      duration  codec_mpeg4  codec_h264  codec_vp8  codec_flv  width  height  \
0  130.35667           1           0           0           0    176    144
```

1	130.35667	1	0	0	0	176	144
2	130.35667	1	0	0	0	176	144
3	130.35667	1	0	0	0	176	144
4	130.35667	1	0	0	0	176	144

	bitrate	framerate	i	...	o_codec_vp8	o_codec_flv	o_codec_h264	\
0	54590	12.0	27	...	0	0	0	
1	54590	12.0	27	...	0	0	0	
2	54590	12.0	27	...	0	0	0	
3	54590	12.0	27	...	0	0	0	
4	54590	12.0	27	...	0	0	0	

	o_bitrate	o_framerate	o_width	o_height	utime	pixels	o_pixels
0	56000	12.0	176	144	0.612	25344	25344
1	56000	12.0	320	240	0.980	25344	76800
2	56000	12.0	480	360	1.216	25344	172800
3	56000	12.0	640	480	1.692	25344	307200
4	56000	12.0	1280	720	3.456	25344	921600

[5 rows x 27 columns]

```
[34]: data_5.info() # teraz wszystkie pola są numeryczne
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 68759 entries, 0 to 68783
Data columns (total 27 columns):
duration      68759 non-null float64
codec_mpeg4   68759 non-null int64
codec_h264    68759 non-null int64
codec_vp8     68759 non-null int64
codec_flv     68759 non-null int64
width         68759 non-null int64
height        68759 non-null int64
bitrate       68759 non-null int64
framerate     68759 non-null float64
i             68759 non-null int64
p             68759 non-null int64
b             68759 non-null int64
frames        68759 non-null int64
i_size        68759 non-null int64
p_size        68759 non-null int64
size          68759 non-null int64
o_codec_mpeg4 68759 non-null int64
o_codec_vp8   68759 non-null int64
o_codec_flv   68759 non-null int64
o_codec_h264  68759 non-null int64
o_bitrate     68759 non-null int64
```

```

o_framerate      68759 non-null float64
o_width          68759 non-null int64
o_height         68759 non-null int64
utime            68759 non-null float64
pixels           68759 non-null int64
o_pixels         68759 non-null int64
dtypes: float64(4), int64(23)
memory usage: 14.7 MB

```

0.8.1 Podział na 'target' i 'features'

```
[35]: target = data_5['utime']
      features = data_5.drop('utime', axis=1)
```

```
[36]: print(features.shape)
      features.head(2)
```

```
(68759, 26)
```

```
[36]:      duration  codec_mpeg4  codec_h264  codec_vp8  codec_flv  width  height  \
0   130.35667           1           0           0           0    176    144
1   130.35667           1           0           0           0    176    144
```

```

      bitrate  framerate  i  ...  o_codec_mpeg4  o_codec_vp8  o_codec_flv  \
0     54590        12.0  27  ...              1           0           0
1     54590        12.0  27  ...              1           0           0

```

```

      o_codec_h264  o_bitrate  o_framerate  o_width  o_height  pixels  o_pixels
0                0     56000         12.0     176     144   25344   25344
1                0     56000         12.0     320     240   25344   76800

```

```
[2 rows x 26 columns]
```

```
[37]: print(target.shape)
      target.head(2)
```

```
(68759,)
```

```
[37]: 0    0.612
      1    0.980
      Name: utime, dtype: float64
```

0.8.2 Skalowanie danych

```
[38]: from sklearn.preprocessing import MinMaxScaler

      scaler = MinMaxScaler()
```

```
features_scaled = scaler.fit_transform(features)
features_scaled[1]
```

```
[38]: array([5.71257598e-02, 1.00000000e+00, 0.00000000e+00, 0.00000000e+00,
          0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 6.06371427e-03,
          1.48820418e-01, 1.00959112e-02, 3.96760662e-02, 0.00000000e+00,
          3.96738187e-02, 5.81774952e-04, 1.02892969e-03, 8.65023435e-04,
          1.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,
          0.00000000e+00, 0.00000000e+00, 8.25688073e-02, 1.02564103e-01,
          0.00000000e+00, 2.51218598e-02])
```

```
[39]: features_scaled.shape
```

```
[39]: (68759, 26)
```

0.9 Trenowanie modeli

```
[40]: from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_squared_error
      from sklearn.model_selection import StratifiedKFold
      from sklearn.metrics import mean_squared_error
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import PolynomialFeatures
      from sklearn.linear_model import ElasticNet
      from sklearn.model_selection import GridSearchCV
```

0.9.1 Podział na dane trenujące i testujące

```
[41]: X_train, X_test, y_train, y_test = train_test_split(features_scaled,
                                                         target,
                                                         test_size=0.3,
                                                         random_state=0)
```

0.9.2 Podstawowy model regresji liniowej

```
[42]: lr = LinearRegression()
      lr.fit(X_train, y_train)

      print("Mean squared error of a linear model: %.2f" %
            mean_squared_error(y_test, lr.predict(X_test)))
      score = lr.score(X_test, y_test) #r2_score
      print("Linear Regression R2 score: %.2f" % score)
```

Mean squared error of a linear model: 124.55

Linear Regression R2 score: 0.53

0.9.3 Wyszukiwanie optymalnego modelu z wykorzystaniem Polynomial Features, ElasticNet oraz GridSearch

```
[43]: parameters=[{
    'alpha':[0.1, 0.2, 0.5, 0.9],
    'l1_ratio':[0.1, 0.2, 0.5, 0.9],
    }
]

for pf_level in range(1,3):
    print(f"polynomial features level {pf_level}:")

    pf = PolynomialFeatures(pf_level)
    train_poly = pf.fit_transform(X_train)
    test_poly = pf.fit_transform(X_test)

    model = GridSearchCV( ElasticNet(), parameters, cv=5 )
    model.fit(train_poly, y_train)

    print(f" {model.best_estimator_}")
    print("    Mean squared error: %.2f" %
mean_squared_error(y_test, model.predict(test_poly)))
    score = model.score(test_poly, y_test) #r2_score
    print("    R2 score: %.2f" % score)
```

polynomial features level 1:

```
ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.9,
           max_iter=1000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Mean squared error: 127.97
R2 score: 0.52
```

polynomial features level 2:

```
ElasticNet(alpha=0.1, copy_X=True, fit_intercept=True, l1_ratio=0.9,
           max_iter=1000, normalize=False, positive=False, precompute=False,
           random_state=None, selection='cyclic', tol=0.0001, warm_start=False)
Mean squared error: 57.47
R2 score: 0.78
```